

НАЦИОНАЛЬНЫЙ ИССЛЕДОВАТЕЛЬСКИЙ УНИВЕРСИТЕТ ИТМО

Факультет ПИиКТ

Системы искусственного интеллекта

Лабораторная работа № 5

Выполнил студент

Набокова Алиса Владиславовна

Группа № Р3320

Преподаватель: Королёва Юлия Александровна

г. Санкт-Петербург

2025

Задание:

Необходимо реализовать простую нейронную сеть(не обязательно сверточную(!)) и обучить ее на датасете

- Каждый слой прописывается руками, перекладыванием матричек/тензоров из состояния А в состояние Б
- Возможные слои/классы, которые вам понадобятся для успешного закрытия лабы: Linear/Dense, Dropout, Conv2d, MaxPool2D, ReLU, Softmax, Flatten, Adam.
- В классе не должно быть волшебных истинно торчовских методов
- Выводим графики нескольких моделей, с экспериментами (разный порядок слоев, разные входные/выходные параметры), выбираем свою лучшую модельку
- dataset <https://www.kaggle.com/datasets/hojjatk/mnist-dataset>

На листингах 1-4 представлены реализации основных слоёв нейронной сети: полносвязный слой, функция активации ReLU, слой Softmax и слой преобразования входных данных Flatten. Каждый слой реализует прямое и обратное распространение ошибки

Листинг 1. Реализация полносвязного слоя (Linear)

```
class Linear:  
    def __init__(self, in_features, out_features):  
        self.W = torch.randn(in_features, out_features) * 0.01  
        self.b = torch.zeros(out_features)  
  
    def forward(self, x):  
        self.x = x  
        return x @ self.W + self.b  
  
    def backward(self, grad):
```

```

    self.dW = self.x.T @ grad
    self.db = grad.sum(dim=0)
    return grad @ self.W.T

```

Листинг 2. Реализация функции активации ReLU

```

class ReLU:
    def forward(self, x):
        self.mask = x > 0
        return x * self.mask

    def backward(self, grad):
        return grad * self.mask

```

Листинг 3. Реализация слоя Softmax

```

class Softmax:
    def forward(self, x):
        exp = torch.exp(x - x.max(dim=1, keepdim=True).values)
        self.out = exp / exp.sum(dim=1, keepdim=True)
        return self.out

    def backward(self, grad):
        return grad

```

Листинг 4. Реализация слоя Flatten

```

class Flatten:
    def forward(self, x):

```

```

    self.shape = x.shape

    return x.view(x.shape[0], -1)

def backward(self, grad):
    return grad.view(self.shape)

```

На листингах 5-6 представлены реализация функции потерь кросс-энтропии и вспомогательная функция вычисления точности классификации

Листинг 5. Реализация функции потерь CrossEntropyLoss

```

class CrossEntropyLoss:

    def forward(self, probs, targets):
        self.probs = probs
        self.targets = targets
        return -torch.log(probs[range(len(targets)), targets]).mean()

    def backward(self):
        grad = self.probs.clone()
        grad[range(len(self.targets)), self.targets] -= 1
        return grad / len(self.targets)

```

Листинг 6. Функция вычисления точности классификации

```

def accuracy_from_probs(probs, labels):
    preds = torch.argmax(probs, dim=1)
    return (preds == labels).float().mean().item()

```

На листингах 7-8 представлены реализация слоя Dropout для регуляризации модели и алгоритма оптимизации Adam

Листинг 7. Реализация слоя Dropout

```
class Dropout:

    def __init__(self, p=0.5):
        self.p = p

    def forward(self, x, train=True):
        if train:
            self.mask = (torch.rand_like(x) > self.p)
            return x * self.mask
        return x

    def backward(self, grad):
        return grad * self.mask
```

Листинг 8. Реализация алгоритма оптимизации Adam

```
class Adam:

    def __init__(self, params, lr=1e-3, betal=0.9, beta2=0.999):
        self.params = params
        self.lr = lr
        self.m = [torch.zeros_like(p) for p in params]
```

```

    self.v = [torch.zeros_like(p) for p in params]

    self.t = 0


def step(self, grads):
    self.t += 1

    for i, (p, g) in enumerate(zip(self.params, grads)):
        self.m[i] = 0.9 * self.m[i] + 0.1 * g
        self.v[i] = 0.999 * self.v[i] + 0.001 * (g ** 2)

        p -= self.lr * self.m[i] / (torch.sqrt(self.v[i]) + 1e-8)

```

На листингах 9-10 представлен процесс инициализации нейронной сети и цикл её обучения с вычислением функции потерь и точности на обучающей и валидационной выборках

Листинг 9. Инициализация архитектуры нейронной сети

```

flatten = Flatten()

fc1 = Linear(784, 256)
relu1 = ReLU()

dropout = Dropout(p=0.2)

fc2 = Linear(256, 128)
relu2 = ReLU()

fc3 = Linear(128, 10)

softmax = Softmax()

loss_fn = CrossEntropyLoss()

params = [fc1.W, fc1.b, fc2.W, fc2.b, fc3.W, fc3.b]
optimizer = Adam(params, lr=1e-3)

```

Листинг 10. Цикл обучения и валидации нейронной сети

```
train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []

epochs = 15

for epoch in range(epochs):
    #train
    total_train_loss = 0
    total_train_acc = 0
    n_batches = 0
    for images, labels in train_loader:

        x = flatten.forward(images)
        x = fc1.forward(x)
        x = relu1.forward(x)
        x = dropout.forward(x, train=True)
        x = fc2.forward(x)
        x = relu2.forward(x)
        x = fc3.forward(x)
        probs = softmax.forward(x)

        loss = loss_fn.forward(probs, labels)
        total_train_loss += loss.item()

    train_losses.append(total_train_loss / n_batches)
    train_accuracies.append(total_train_acc / n_batches)

    #val
    total_val_loss = 0
    total_val_acc = 0
    n_batches = 0
    for images, labels in val_loader:

        x = flatten.forward(images)
        x = fc1.forward(x)
        x = relu1.forward(x)
        x = dropout.forward(x, train=False)
        x = fc2.forward(x)
        x = relu2.forward(x)
        x = fc3.forward(x)
        probs = softmax.forward(x)

        loss = loss_fn.forward(probs, labels)
        total_val_loss += loss.item()

    val_losses.append(total_val_loss / n_batches)
    val_accuracies.append(total_val_acc / n_batches)
```

```

total_train_acc += accuracy_from_probs(probs, labels)

n_batches += 1

grad = loss_fn.backward()

grad = softmax.backward(grad)

grad = fc3.backward(grad)

grad = relu2.backward(grad)

grad = fc2.backward(grad)

grad = dropout.backward(grad)

grad = relu1.backward(grad)

grad = fc1.backward(grad)

grads = [fc1.dW, fc1.db, fc2.dW, fc2.db, fc3.dW, fc3.db]

optimizer.step(grads)

train_losses.append(total_train_loss / n_batches)

train_accuracies.append(total_train_acc / n_batches)

#validation

total_val_loss = 0

total_val_acc = 0

n_val_batches = 0

for images, labels in val_loader:

    x = flatten.forward(images)

    x = fc1.forward(x)

```

```

        x = relu1.forward(x)

        x = fc2.forward(x)

        x = relu2.forward(x)

        x = fc3.forward(x)

        probs = softmax.forward(x)

        loss = loss_fn.forward(probs, labels)

        total_val_loss += loss.item()

        total_val_acc += accuracy_from_probs(probs, labels)

        n_val_batches += 1

    val_losses.append(total_val_loss / n_val_batches)

    val_accuracies.append(total_val_acc / n_val_batches)

print(
    f"Epoch {epoch+1}: "
    f"train_loss={train_losses[-1]:.4f}, "
    f"train_acc={train_accuracies[-1]*100:.2f}%, "
    f"val_loss={val_losses[-1]:.4f}, "
    f"val_acc={val_accuracies[-1]*100:.2f}%""
)

```

На листингах 11-14 представлены функции оценки качества модели на тестовой выборке, сохранения и загрузки параметров модели, а также пример применения обученной модели для распознавания рукописного изображения

Листинг 11. Оценка точности на тестовой выборке

```
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

axes[0].plot(train_losses, label="train loss")
axes[0].plot(val_losses, label="val loss")
axes[0].set_xlabel("Epoch")
axes[0].set_ylabel("Loss")
axes[0].legend()
axes[0].grid(True)
axes[0].set_title("Loss")

axes[1].plot(train_accuracies, label="train accuracy")
axes[1].plot(val_accuracies, label="val accuracy")
axes[1].set_xlabel("Epoch")
axes[1].set_ylabel("Accuracy")
axes[1].legend()
axes[1].grid(True)
axes[1].set_title("Accuracy")

plt.tight_layout()
plt.show()
```

Листинг 12. Оценка качества модели на тестовой выборке

```
def accuracy(loader):

    correct = 0

    total = 0

    for images, labels in loader:

        x = flatten.forward(images)

        x = fc1.forward(x)

        x = relu1.forward(x)

        x = fc2.forward(x)

        x = relu2.forward(x)

        x = fc3.forward(x)

        probs = softmax.forward(x)

        preds = torch.argmax(probs, dim=1)

        correct += (preds == labels).sum().item()

        total += labels.size(0)

    return correct / total

print("Test accuracy:", accuracy(test_loader))
```

Листинг 13. Сохранение модели

```
def save_model(filename="mymodel.pth"):

    state = {

        "fc1.W": fc1.W,

        "fc1.b": fc1.b,

        "fc2.W": fc2.W,

        "fc2.b": fc2.b,

        "fc3.W": fc3.W,
```

```

    "fc3.b": fc3.b,
}

torch.save(state, filename)

```

Листинг 14. Загрузка модели

```

def load_model(filename="mymodel.pth"):

    state = torch.load(filename)

    fc1.W = state["fc1.W"]

    fc1.b = state["fc1.b"]

    fc2.W = state["fc2.W"]

    fc2.b = state["fc2.b"]

    fc3.W = state["fc3.W"]

    fc3.b = state["fc3.b"]

```

Листинг 14. Распознавание изображения

```

load_model("/content/mymodel_with_1hidden_15ep_reg.pth")

img = Image.open("/content/0.png").convert("L")

img = img.resize((28, 28))

img = PIL.ImageOps.invert(img)

x = transform(img).unsqueeze(0)

x_flat = flatten.forward(x)

x1 = fc1.forward(x_flat)

x1 = relu1.forward(x1)

x2 = fc2.forward(x1)

x2 = relu2.forward(x2)

x3 = fc3.forward(x2)

```

```
probs = softmax.forward(x3)

pred = torch.argmax(probs, dim=1).item()

print("Predicted digit:", pred)
```

Эксперименты:

Model 1

Linear(784, 128)

relu

Linear(128, 10)

softmax

epochs: 5, lr=0.001

оценка модели:

Epoch 1:	train_loss=0.3176,	train_acc=90.33%,	val_loss=0.2385,
val_acc=93.23%			
Epoch 2:	train_loss=0.1711,	train_acc=94.92%,	val_loss=0.1872,
val_acc=94.43%			
Epoch 3:	train_loss=0.1362,	train_acc=95.86%,	val_loss=0.1507,
val_acc=95.88%			
Epoch 4:	train_loss=0.1166,	train_acc=96.46%,	val_loss=0.1555,
val_acc=95.83%			
Epoch 5:	train_loss=0.1026,	train_acc=96.93%,	val_loss=0.1458,
val_acc=95.88%			




Рисунок 1. Потери и точности для модели 1

Test accuracy: 0.9687

Model 2

Linear(784, 128)

relu

dropout(p=0.2)

Linear(128, 10)

softmax

epochs: 5, lr=0.001

оценка модели:

```

Epoch    1:    train_loss=0.4008,      train_acc=87.49%,      val_loss=0.2447,
val_acc=92.92%

Epoch    2:    train_loss=0.2510,      train_acc=92.46%,      val_loss=0.2220,
val_acc=93.82%

Epoch    3:    train_loss=0.2144,      train_acc=93.56%,      val_loss=0.1878,
val_acc=94.70%

Epoch    4:    train_loss=0.1941,      train_acc=94.06%,      val_loss=0.1774,
val_acc=95.13%
  
```

```
Epoch      5:    train_loss=0.1837,    train_acc=94.43%,    val_loss=0.1623,  
val_acc=95.45%
```




Рисунок 2. Потери и точности для модели 2

Test accuracy: 0.9489

Model 3

Linear(784, 256)

relu

dropout(p=0.2)

Linear(256, 10)

softmax

epochs: 5, lr=0.001

оценка модели:

```
Epoch      1:    train_loss=0.3363,    train_acc=89.47%,    val_loss=0.1993,  
val_acc=94.37%
```

```
Epoch      2:    train_loss=0.1897,    train_acc=94.30%,    val_loss=0.1610,  
val_acc=95.50%
```

```

Epoch    3:    train_loss=0.1567,    train_acc=95.14%,    val_loss=0.1435,
val_acc=95.88%

Epoch    4:    train_loss=0.1379,    train_acc=95.77%,    val_loss=0.1504,
val_acc=95.87%

Epoch    5:    train_loss=0.1247,    train_acc=96.13%,    val_loss=0.1366,
val_acc=96.38%

```




Рисунок 3. Потери и точности для модели 3

Test accuracy: 0.9627

Model 4

Linear(784, 256)

relu

Linear(256, 128)

relu

Linear(128, 10)

softmax

epochs: 5, lr=0.001

оценка модели:

```

Epoch 1: train_loss=0.3197, train_acc=89.83%, val_loss=0.1885,
val_acc=94.60%

Epoch 2: train_loss=0.1465, train_acc=95.43%, val_loss=0.1689,
val_acc=94.87%

Epoch 3: train_loss=0.1139, train_acc=96.55%, val_loss=0.1345,
val_acc=96.13%

Epoch 4: train_loss=0.0947, train_acc=97.03%, val_loss=0.1315,
val_acc=96.35%

Epoch 5: train_loss=0.0804, train_acc=97.41%, val_loss=0.1044,
val_acc=96.93%

```




Рисунок 4. Потери и точности для модели 4

Test accuracy: 0.9714

Model 5

Linear(784, 256)

relu

dropout(p=0.25)

Linear(256, 128)

relu

Linear(128, 10)

softmax

epochs: 10, lr=0.001

оценка модели:

Epoch 1:	train_loss=0.4001,	train_acc=87.20%,	val_loss=0.2211,
val_acc=93.83%			
Epoch 2:	train_loss=0.2145,	train_acc=93.40%,	val_loss=0.1871,
val_acc=94.98%			
Epoch 3:	train_loss=0.1810,	train_acc=94.27%,	val_loss=0.1795,
val_acc=95.27%			
Epoch 4:	train_loss=0.1631,	train_acc=94.91%,	val_loss=0.1641,
val_acc=95.85%			
Epoch 5:	train_loss=0.1478,	train_acc=95.42%,	val_loss=0.1392,
val_acc=96.48%			
Epoch 6:	train_loss=0.1394,	train_acc=95.57%,	val_loss=0.1398,
val_acc=96.60%			
Epoch 7:	train_loss=0.1305,	train_acc=95.90%,	val_loss=0.1299,
val_acc=96.55%			
Epoch 8:	train_loss=0.1238,	train_acc=96.05%,	val_loss=0.1351,
val_acc=96.80%			
Epoch 9:	train_loss=0.1159,	train_acc=96.30%,	val_loss=0.1404,
val_acc=96.75%			
Epoch 10:	train_loss=0.1163,	train_acc=96.28%,	val_loss=0.1313,
val_acc=96.82%			




Рисунок 5. Потери и точности для модели 5

Test accuracy: 0.9696

Model 6

Linear(784, 512)

relu

dropout(p=0.2)

Linear(512, 256)

relu

Linear(256, 10)

softmax

epochs: 15, lr=0.001

оценка модели:

Epoch 1:	train_loss=0.3692,	train_acc=88.34%,	val_loss=0.2043,
val_acc=94.28%			
Epoch 2:	train_loss=0.1904,	train_acc=94.05%,	val_loss=0.1541,
val_acc=95.40%			
Epoch 3:	train_loss=0.1593,	train_acc=95.09%,	val_loss=0.1559,
val_acc=95.82%			
Epoch 4:	train_loss=0.1405,	train_acc=95.56%,	val_loss=0.1506,
val_acc=95.78%			
Epoch 5:	train_loss=0.1296,	train_acc=96.00%,	val_loss=0.1418,
val_acc=96.20%			
Epoch 6:	train_loss=0.1174,	train_acc=96.26%,	val_loss=0.1290,
val_acc=96.53%			
Epoch 7:	train_loss=0.1102,	train_acc=96.52%,	val_loss=0.1401,
val_acc=96.75%			
Epoch 8:	train_loss=0.1072,	train_acc=96.61%,	val_loss=0.1226,
val_acc=96.78%			
Epoch 9:	train_loss=0.0975,	train_acc=96.91%,	val_loss=0.1168,
val_acc=97.07%			

```

Epoch    10:    train_loss=0.0955,    train_acc=96.87%,    val_loss=0.1076,
val_acc=97.40%

Epoch    11:    train_loss=0.0911,    train_acc=97.05%,    val_loss=0.1130,
val_acc=97.00%

Epoch    12:    train_loss=0.0847,    train_acc=97.25%,    val_loss=0.1052,
val_acc=97.35%

Epoch    13:    train_loss=0.0827,    train_acc=97.27%,    val_loss=0.1201,
val_acc=97.17%

Epoch    14:    train_loss=0.0846,    train_acc=97.23%,    val_loss=0.1147,
val_acc=97.28%

Epoch    15:    train_loss=0.0784,    train_acc=97.50%,    val_loss=0.1102,
val_acc=97.00%

```




Рисунок 6. Потери и точности для модели 6

Test accuracy: 0.9749

Model 7

Linear(784, 512)

relu

dropout(p=0.2)

Linear(512, 256)

relu

Linear(256, 10)

softmax

epochs: 30, lr=0.001

оценка модели:

Epoch 1:	train_loss=0.3599,	train_acc=88.46%,	val_loss=0.1891,
	val_acc=94.40%		
Epoch 2:	train_loss=0.1850,	train_acc=94.38%,	val_loss=0.1757,
	val_acc=95.17%		
Epoch 3:	train_loss=0.1480,	train_acc=95.33%,	val_loss=0.1434,
	val_acc=95.95%		
Epoch 4:	train_loss=0.1319,	train_acc=95.87%,	val_loss=0.1594,
	val_acc=95.92%		
Epoch 5:	train_loss=0.1222,	train_acc=96.15%,	val_loss=0.1254,
	val_acc=96.90%		
Epoch 6:	train_loss=0.1086,	train_acc=96.51%,	val_loss=0.1319,
	val_acc=96.57%		
Epoch 7:	train_loss=0.0994,	train_acc=96.80%,	val_loss=0.1151,
	val_acc=97.05%		
Epoch 8:	train_loss=0.0943,	train_acc=96.97%,	val_loss=0.1249,
	val_acc=96.87%		
Epoch 9:	train_loss=0.0871,	train_acc=97.14%,	val_loss=0.1168,
	val_acc=97.20%		
Epoch 10:	train_loss=0.0839,	train_acc=97.30%,	val_loss=0.1317,
	val_acc=96.87%		
Epoch 11:	train_loss=0.0790,	train_acc=97.44%,	val_loss=0.1092,
	val_acc=97.50%		

```
    Epoch    12:    train_loss=0.0767,    train_acc=97.43%,    val_loss=0.1283,  
val_acc=97.23%  
  
    Epoch    13:    train_loss=0.0711,    train_acc=97.67%,    val_loss=0.1117,  
val_acc=97.52%  
  
    Epoch    14:    train_loss=0.0714,    train_acc=97.67%,    val_loss=0.1288,  
val_acc=97.10%  
  
    Epoch    15:    train_loss=0.0698,    train_acc=97.69%,    val_loss=0.1057,  
val_acc=97.73%  
  
    Epoch    16:    train_loss=0.0663,    train_acc=97.84%,    val_loss=0.1255,  
val_acc=97.20%  
  
    Epoch    17:    train_loss=0.0637,    train_acc=97.88%,    val_loss=0.1130,  
val_acc=97.62%  
  
    Epoch    18:    train_loss=0.0609,    train_acc=97.96%,    val_loss=0.1078,  
val_acc=97.52%  
  
    Epoch    19:    train_loss=0.0613,    train_acc=97.91%,    val_loss=0.1172,  
val_acc=97.52%  
  
    Epoch    20:    train_loss=0.0589,    train_acc=98.05%,    val_loss=0.1155,  
val_acc=97.47%  
  
    Epoch    21:    train_loss=0.0612,    train_acc=97.92%,    val_loss=0.1166,  
val_acc=97.57%  
  
    Epoch    22:    train_loss=0.0538,    train_acc=98.20%,    val_loss=0.1297,  
val_acc=97.20%  
  
    Epoch    23:    train_loss=0.0559,    train_acc=98.17%,    val_loss=0.1111,  
val_acc=97.55%  
  
    Epoch    24:    train_loss=0.0552,    train_acc=98.16%,    val_loss=0.1144,  
val_acc=97.43%  
  
    Epoch    25:    train_loss=0.0549,    train_acc=98.19%,    val_loss=0.1193,  
val_acc=97.45%  
  
    Epoch    26:    train_loss=0.0519,    train_acc=98.29%,    val_loss=0.1139,  
val_acc=97.47%  
  
    Epoch    27:    train_loss=0.0509,    train_acc=98.31%,    val_loss=0.1146,  
val_acc=97.57%
```

```

Epoch    28:    train_loss=0.0501,    train_acc=98.36%,    val_loss=0.1217,
val_acc=97.50%

Epoch    29:    train_loss=0.0515,    train_acc=98.28%,    val_loss=0.1237,
val_acc=97.38%

Epoch    30:    train_loss=0.0490,    train_acc=98.37%,    val_loss=0.1445,
val_acc=97.07%

```




Рисунок 7. Потери и точности для модели 7

Test accuracy: 0.9754

Model 8

Linear(784, 256)

relu

dropout(p=0.2)

Linear(256, 128)

relu

Linear(128, 10)

softmax

epochs: 15, lr=0.001

оценка модели:

```
Epoch    1:    train_loss=0.3823,    train_acc=87.75%,    val_loss=0.2166,
val_acc=93.60%  
  
Epoch    2:    train_loss=0.2063,    train_acc=93.62%,    val_loss=0.1815,
val_acc=94.87%  
  
Epoch    3:    train_loss=0.1722,    train_acc=94.60%,    val_loss=0.1521,
val_acc=95.75%  
  
Epoch    4:    train_loss=0.1521,    train_acc=95.21%,    val_loss=0.1547,
val_acc=95.85%  
  
Epoch    5:    train_loss=0.1414,    train_acc=95.56%,    val_loss=0.1289,
val_acc=96.40%  
  
Epoch    6:    train_loss=0.1305,    train_acc=95.89%,    val_loss=0.1444,
val_acc=95.72%  
  
Epoch    7:    train_loss=0.1209,    train_acc=96.13%,    val_loss=0.1248,
val_acc=96.60%  
  
Epoch    8:    train_loss=0.1144,    train_acc=96.41%,    val_loss=0.1193,
val_acc=96.83%  
  
Epoch    9:    train_loss=0.1084,    train_acc=96.51%,    val_loss=0.1279,
val_acc=96.93%  
  
Epoch   10:    train_loss=0.1029,    train_acc=96.70%,    val_loss=0.1299,
val_acc=96.67%  
  
Epoch   11:    train_loss=0.0947,    train_acc=96.90%,    val_loss=0.1329,
val_acc=96.73%  
  
Epoch   12:    train_loss=0.0961,    train_acc=96.88%,    val_loss=0.1152,
val_acc=97.20%
```

```

Epoch    13:    train_loss=0.0939,    train_acc=96.99%,    val_loss=0.1257,
val_acc=96.73%

Epoch    14:    train_loss=0.0880,    train_acc=97.16%,    val_loss=0.1196,
val_acc=97.22%

Epoch    15:    train_loss=0.0853,    train_acc=97.20%,    val_loss=0.1238,
val_acc=97.08%

```




Рисунок 8. Потери и точности для модели 8

Test accuracy: 0.9719




Рисунок 9. ROC-кривая для 10 классов выборки

Вывод

В процессе экспериментов было рассмотрено несколько конфигураций архитектуры и гиперпараметров модели, отличающихся числом скрытых слоёв, размерностью скрытых представлений, использованием регуляризации и количеством эпох обучения. По результатам сравнительного анализа лучшей оказалась модель Model 8 со следующей архитектурой:

- входной слой: Linear(784, 256),
- функция активации ReLU,
- слой регуляризации Dropout с вероятностью отключения нейронов $p = 0.2$,
- скрытый слой Linear(256, 128) с функцией активации ReLU,
- выходной слой Linear(128, 10) с функцией Softmax.

Обучение модели проводилось в течение 15 эпох с использованием оптимизатора Adam и шагом обучения $lr = 0.001$

В результате обучения было достигнуто:

- точность на обучающей выборке: 97.20%
- точность на валидационной выборке: 97.22%
- точность классификации на тестовой выборке: 97.19%

Использование слоя Dropout позволило снизить эффект переобучения и обеспечить более стабильное поведение модели на валидационных и тестовых данных по сравнению с моделями без регуляризации. При увеличении количества эпох без Dropout наблюдалось переобучение, выражющееся в росте ошибки на валидационной выборке.

Проведённые эксперименты показали, что увеличение сложности модели сверх двух скрытых слоёв не приводит к существенному росту качества, но повышает риск переобучения. Таким образом, выбранная архитектура является компромиссом между точностью, устойчивостью и вычислительной сложностью.

Разработанная модель успешно решает задачу распознавания рукописных цифр и может применяться для классификации изображений